

NLP_Project

April 21, 2023

DISASTER TWEETS PREDICTION USING NATURAL LANGUAGE PROCESSING

TEAM MEMBERS:

T.S.S. ABINANDHAN KUMAR | 19MIA1062
NIRANJAN J | 19MIA1003
ALAGARSAMY N | 19MIA1082

```
[3]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

1 1. Importing the necessary libraries

```
[4]: import numpy as np
import pandas as pd

# text processing libraries
import re
import string
import nltk
from nltk.corpus import stopwords

# XGBoost
import xgboost as xgb
from xgboost import XGBClassifier

# sklearn
from sklearn import model_selection
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import f1_score
```

```

from sklearn import preprocessing, decomposition, model_selection, metrics, u
↳pipeline
from sklearn.model_selection import u
↳GridSearchCV, StratifiedKFold, RandomizedSearchCV

# matplotlib and seaborn for plotting
import matplotlib.pyplot as plt
import seaborn as sns

# File system management
import os

# Suppress warnings
import warnings
warnings.filterwarnings('ignore')

```

2 2. Reading the datasets

[5]: #Training data
train = pd.read_csv('/content/drive/MyDrive/NLP/train.csv')
train.head()

[5]: id keyword location text \\\n0 1 NaN NaN Our Deeds are the Reason of this #earthquake M...\\n1 4 NaN NaN Forest fire near La Ronge Sask. Canada\\n2 5 NaN NaN All residents asked to 'shelter in place' are ...\\n3 6 NaN NaN 13,000 people receive #wildfires evacuation or...\\n4 7 NaN NaN Just got sent this photo from Ruby #Alaska as ...

target\\n0 1\\n1 1\\n2 1\\n3 1\\n4 1

[6]: test = pd.read_csv('/content/drive/MyDrive/NLP/test.csv')
test.head()

[6]: id keyword location text\\n0 0 NaN NaN Just happened a terrible car crash\\n1 2 NaN NaN Heard about #earthquake is different cities, s...\\n2 3 NaN NaN there is a forest fire at spot pond, geese are...\\n3 9 NaN NaN Apocalypse lighting. #Spokane #wildfires\\n4 11 NaN NaN Typhoon Soudelor kills 28 in China and Taiwan

3 3. Basic EDA

3.1 Missing values

```
[7]: train.isnull().sum()
```

```
[7]: id          0  
keyword      61  
location    2533  
text         0  
target        0  
dtype: int64
```

```
[8]: test.isnull().sum()
```

```
[8]: id          0  
keyword      26  
location    1105  
text         0  
dtype: int64
```

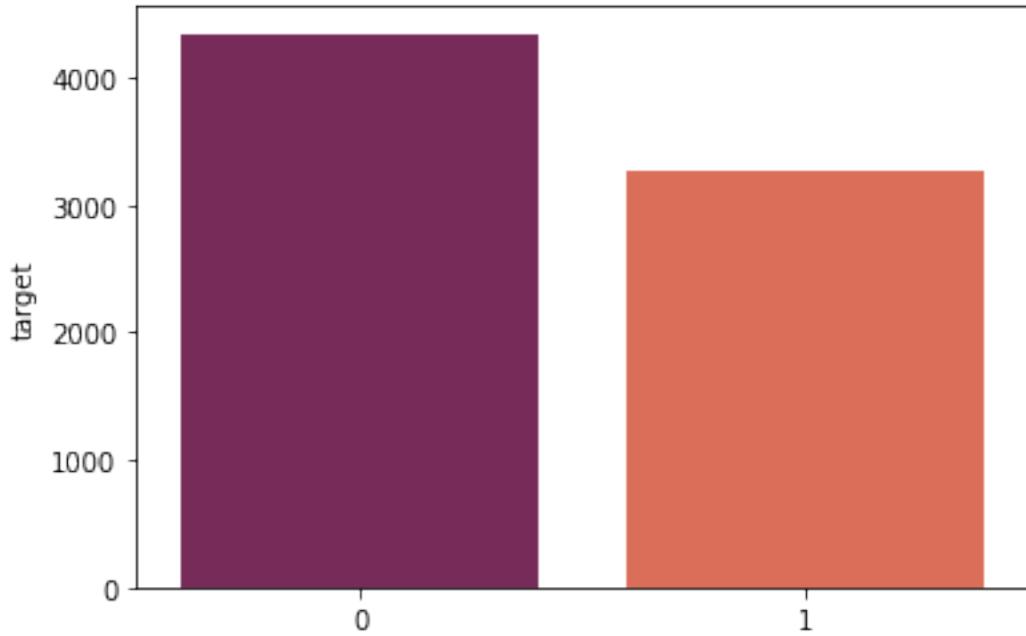
3.2 Exploring the Target Column

```
[9]: train['target'].value_counts()
```

```
[9]: 0    4342  
1    3271  
Name: target, dtype: int64
```

```
[10]: sns.barplot(train['target'].value_counts().index,train['target'] .  
    ↪value_counts(),palette='rocket')
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5e8c856090>
```



```
[11]: disaster_tweets = train[train['target']==1]['text']
disaster_tweets.values[1]
```

```
[11]: 'Forest fire near La Ronge Sask. Canada'
```

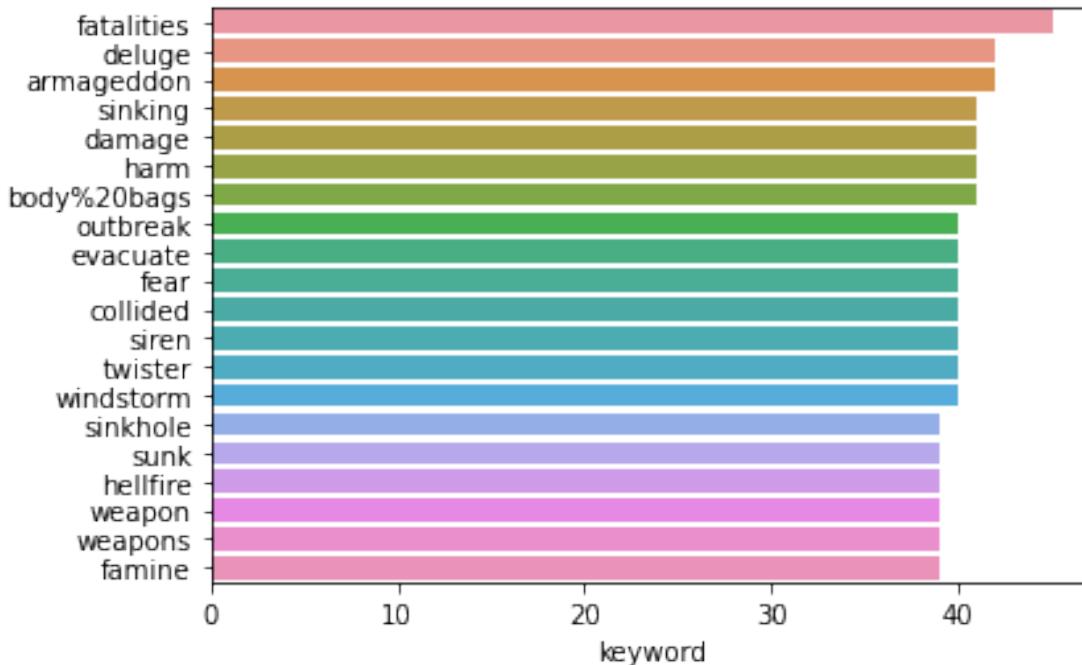
```
[12]: non_disaster_tweets = train[train['target']==0]['text']
non_disaster_tweets.values[1]
```

```
[12]: 'I love fruits'
```

3.3 Exploring the 'keyword' column

```
[13]: sns.barplot(y=train['keyword'].value_counts()[:20].index,x=train['keyword'].
    ↴value_counts()[:20],
    orient='h')
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5e8c73c110>
```



```
[14]: train.loc[train['text'].str.contains('disaster', na=False, case=False)].target.
      ↴value_counts()
```

```
[14]: 1    102
      0     40
      Name: target, dtype: int64
```

3.4 Exploring the ‘location’ column

```
[15]: train['location'].replace({'United States':'USA',
                               'New York':'USA',
                               "London":'UK',
                               "Los Angeles, CA":'USA',
                               "Washington, D.C.":'USA',
                               "California":'USA',
                               "Chicago, IL":'USA',
                               "Chicago":'USA',
                               "New York, NY":'USA',
                               "California, USA":'USA',
                               "Florida":'USA',
                               "Nigeria":'Africa',
                               "Kenya":'Africa',
                               "Everywhere":'Worldwide',
                               "San Francisco":'USA',
                               "Florida":'USA',
```

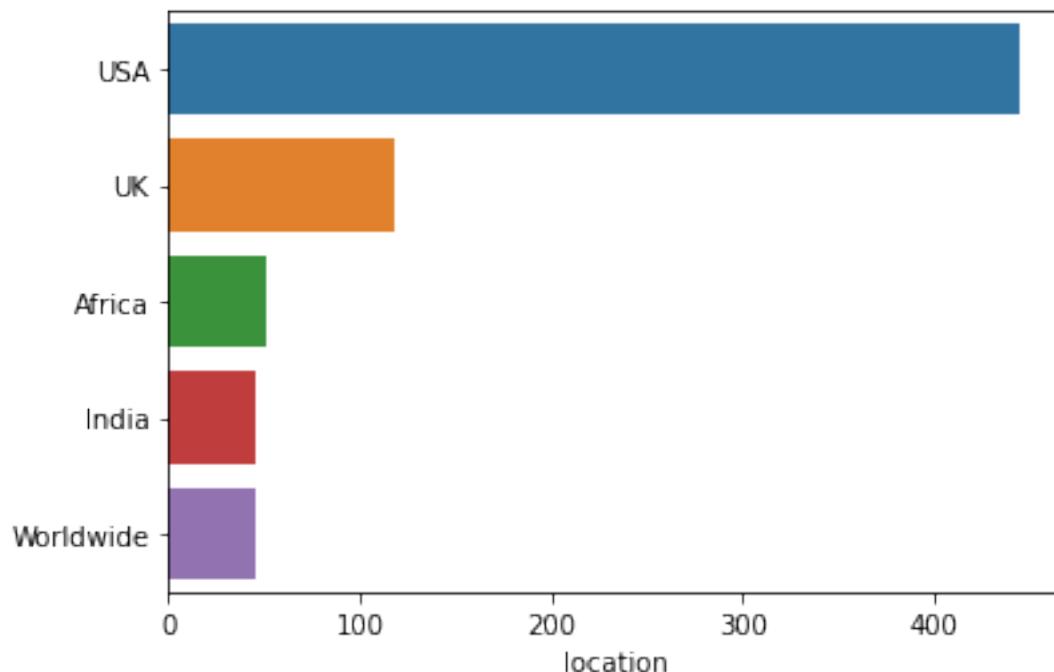
```

    "United Kingdom":'UK',
    "Los Angeles":'USA',
    "Toronto":'Canada',
    "San Francisco, CA":'USA',
    "NYC":'USA',
    "Seattle":'USA',
    "Earth":'Worldwide',
    "Ireland":'UK',
    "London, England":'UK',
    "New York City":'USA',
    "Texas":'USA',
    "London, UK":'UK',
    "Atlanta, GA":'USA',
    "Mumbai":'India'},inplace=True)

sns.barplot(y=train['location'].value_counts()[:5].index,x=train['location'].
            value_counts()[:5],
            orient='h')

```

[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5e8c749f90>



[16]: train['text'][:5]

```
[16]: 0    Our Deeds are the Reason of this #earthquake M...
       1          Forest fire near La Ronge Sask. Canada
       2    All residents asked to 'shelter in place' are ...
       3    13,000 people receive #wildfires evacuation or...
       4    Just got sent this photo from Ruby #Alaska as ...
Name: text, dtype: object
```

TEXT PREPROCESSING

```
[17]: def clean_text(text):
        text = text.lower()
        text = re.sub('\[.*?\]', '', text)
        text = re.sub('https://\S+|www\.\S+', '', text)
        text = re.sub('<.*?>+', '', text)
        text = re.sub('[%s]' % re.escape(string.punctuation), '', text)
        text = re.sub('\n', '', text)
        text = re.sub('\w*\d\w*', '', text)
        return text

train['text'] = train['text'].apply(lambda x: clean_text(x))
test['text'] = test['text'].apply(lambda x: clean_text(x))

train['text'].head()
```

```
[17]: 0    our deeds are the reason of this earthquake ma...
       1          forest fire near la ronge sask canada
       2    all residents asked to shelter in place are be...
       3    people receive wildfires evacuation orders in...
       4    just got sent this photo from ruby alaska as s...
Name: text, dtype: object
```

```
[18]: from wordcloud import WordCloud
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=[26, 8])
wordcloud1 = WordCloud( background_color='white',
                       width=600,
                       height=400).generate(" ".join(disaster_tweets))
ax1.imshow(wordcloud1)
ax1.axis('off')
ax1.set_title('Disaster Tweets', fontsize=40);

wordcloud2 = WordCloud( background_color='white',
                       width=600,
                       height=400).generate(" ".join(non_disaster_tweets))
ax2.imshow(wordcloud2)
ax2.axis('off')
ax2.set_title('Non Disaster Tweets', fontsize=40);
```



3.5 2. Tokenization

```
[19]: text = "Are you coming , aren't you"
tokenizer1 = nltk.tokenize.WhitespaceTokenizer()
tokenizer2 = nltk.tokenize.TreebankWordTokenizer()
tokenizer3 = nltk.tokenize.WordPunctTokenizer()
tokenizer4 = nltk.tokenize.RegexpTokenizer(r'\w+')

print("Example Text: ",text)
print("-----")
print("Tokenization by whitespace:- ",tokenizer1.tokenize(text))
print("Tokenization by words using Treebank Word Tokenizer:- ",tokenizer2.
     tokenize(text))
print("Tokenization by punctuation:- ",tokenizer3.tokenize(text))
print("Tokenization by regular expression:- ",tokenizer4.tokenize(text))
```

Example Text: Are you coming , aren't you

```
Tokenization by whitespace:- ['Are', 'you', 'coming', ',', ',', "aren't", 'you']
Tokenization by words using Treebank Word Tokenizer:- ['Are', 'you', 'coming',
', ', 'are', "n't", 'you']
Tokenization by punctuation:- ['Are', 'you', 'coming', ',', 'aren', "", 't',
'you']
Tokenization by regular expression:- ['Are', 'you', 'coming', 'aren', 't',
'you']
```

```
[20]: tokenizer = nltk.tokenize.RegexpTokenizer(r'\w+')
train['text'] = train['text'].apply(lambda x: tokenizer.tokenize(x))
test['text'] = test['text'].apply(lambda x: tokenizer.tokenize(x))
train['text'].head()
```

```
[20]: 0      [our, deeds, are, the, reason, of, this, earth...
       1      [forest, fire, near, la, ronge, sask, canada]
       2      [all, residents, asked, to, shelter, in, place...
       3      [people, receive, wildfires, evacuation, order...
       4      [just, got, sent, this, photo, from, ruby, ala...
Name: text, dtype: object
```

3. STOPWORDS REMOVAL

```
[21]: import nltk
nltk.download('stopwords')
def remove_stopwords(text):

    words = [w for w in text if w not in stopwords.words('english')]
    return words

train['text'] = train['text'].apply(lambda x : remove_stopwords(x))
test['text'] = test['text'].apply(lambda x : remove_stopwords(x))
train.head()
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]  Unzipping corpora/stopwords.zip.
```

	id	keyword	location	text \
0	1	NaN	NaN	[deeds, reason, earthquake, may, allah, forgiv...
1	4	NaN	NaN	[forest, fire, near, la, ronge, sask, canada]
2	5	NaN	NaN	[residents, asked, shelter, place, notified, o...
3	6	NaN	NaN	[people, receive, wildfires, evacuation, order...
4	7	NaN	NaN	[got, sent, photo, ruby, alaska, smoke, wildfi...

	target
0	1
1	1
2	1
3	1
4	1

3.6 4. Token normalization

Token normalisation means converting different tokens to their base forms. This can be done either by:

- **Stemming** : removing and replacing suffixes to get to the root form of the word, which is called the **stem** for instance cats - cat, wolves - wolv
- **Lemmatization** : Returns the base or dictionary form of a word, which is known as the **lemma**

source

```
[22]: nltk.download('wordnet')
# Stemming and Lemmatization examples
text = "feet cats wolves talked"

tokenizer = nltk.tokenize.TreebankWordTokenizer()
tokens = tokenizer.tokenize(text)

# Stemmer
stemmer = nltk.stem.PorterStemmer()
print("Stemming the sentence: ", " ".join(stemmer.stem(token) for token in tokens))

# Lemmatizer
lemmatizer=nltk.stem.WordNetLemmatizer()
print("Lemmatizing the sentence: ", " ".join(lemmatizer.lemmatize(token) for token in tokens))
```

[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Unzipping corpora/wordnet.zip.
Stemming the sentence: feet cat wolv talk
Lemmatizing the sentence: foot cat wolf talked

```
[23]: def combine_text(list_of_text):
    '''Takes a list of text and combines them into one large chunk of text.'''
    combined_text = ' '.join(list_of_text)
    return combined_text

train['text'] = train['text'].apply(lambda x : combine_text(x))
test['text'] = test['text'].apply(lambda x : combine_text(x))
train['text']
train.head()
```

	id	keyword	location	text
0	1	NaN	NaN	deeds reason earthquake may allah forgive us
1	4	NaN	NaN	forest fire near la ronge sask canada
2	5	NaN	NaN	residents asked shelter place notified officer...
3	6	NaN	NaN	people receive wildfires evacuation orders cal...
4	7	NaN	NaN	got sent photo ruby alaska smoke wildfires pou...

	target
0	1
1	1
2	1
3	1
4	1

```
[24]: def text_preprocessing(text):
    """
    Cleaning and parsing the text.

    """
    tokenizer = nltk.tokenize.RegexpTokenizer(r'\w+')

    nopunc = clean_text(text)
    tokenized_text = tokenizer.tokenize(nopunc)
    remove_stopwords = [w for w in tokenized_text if w not in stopwords.
    ↪words('english')]
    combined_text = ' '.join(remove_stopwords)
    return combined_text
```

```
[25]: count_vectorizer = CountVectorizer()
train_vectors = count_vectorizer.fit_transform(train['text'])
test_vectors = count_vectorizer.transform(test["text"])

print(train_vectors[0].todense())
```

[[0 0 0 ... 0 0 0]]

```
[26]: tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
train_tfidf = tfidf.fit_transform(train['text'])
test_tfidf = tfidf.transform(test["text"])
```

4 6. Building a Text Classification model

4.1 Logistic Regression Classifier

```
[27]: clf = LogisticRegression(C=1.0)
scores = model_selection.cross_val_score(clf, train_vectors, train["target"],
    ↪cv=5, scoring="f1")
scores
```

[27]: array([0.59865255, 0.49611063, 0.57166948, 0.56290774, 0.68789809])

```
[28]: clf.fit(train_vectors, train["target"])
```

```
[28]: LogisticRegression()
```

```
[29]: clf_tfidf = LogisticRegression(C=1.0)
scores = model_selection.cross_val_score(clf_tfidf, train_tfidf,
    ↪train["target"], cv=5, scoring="f1")
scores
```

```
[29]: array([0.57229525, 0.49673203, 0.54277829, 0.46618106, 0.64768683])
```

4.2 Naives Bayes Classifier

```
[30]: clf_NB = MultinomialNB()
scores = model_selection.cross_val_score(clf_NB, train_vectors, □
    ↪train["target"], cv=5, scoring="f1")
scores
```

```
[30]: array([0.63149079, 0.60675773, 0.68575519, 0.64341085, 0.72505092])
```

```
[31]: clf_NB.fit(train_vectors, train["target"])
```

```
[31]: MultinomialNB()
```

```
[32]: clf_NB_TFIDF = MultinomialNB()
scores = model_selection.cross_val_score(clf_NB_TFIDF, train_tfidf, □
    ↪train["target"], cv=5, scoring="f1")
scores
```

```
[32]: array([0.57590597, 0.57092511, 0.61135371, 0.5962963 , 0.7393745 ])
```

```
[33]: clf_NB_TFIDF.fit(train_tfidf, train["target"])
```

```
[33]: MultinomialNB()
```

4.3 XGBoost

```
[34]: import xgboost as xgb
clf_xgb = xgb.XGBClassifier(max_depth=7, n_estimators=200, colsample_bytree=0.8,
                             subsample=0.8, nthread=10, learning_rate=0.1)
scores = model_selection.cross_val_score(clf_xgb, train_vectors, □
    ↪train["target"], cv=5, scoring="f1")
scores
```

```
[34]: array([0.47379913, 0.37379576, 0.43988816, 0.38900634, 0.53142857])
```

```
[35]: import xgboost as xgb
clf_xgb_TFIDF = xgb.XGBClassifier(max_depth=7, n_estimators=200, □
    ↪colsample_bytree=0.8,
                             subsample=0.8, nthread=10, learning_rate=0.1)
scores = model_selection.cross_val_score(clf_xgb_TFIDF, train_tfidf, □
    ↪train["target"], cv=5, scoring="f1")
scores
```

```
[35]: array([0.48947951, 0.34406439, 0.43140965, 0.40084388, 0.53014354])
```